

Game Playing

Game Playing

Why do AI researchers study game playing?

1. It's a good reasoning problem, formal and nontrivial.
2. Direct comparison with humans and other computer programs is easy.

Why study games?

- Interesting, hard problems which require minimal “initial structure”
- Clear criteria for success
- Offer an opportunity to study problems involving {hostile, adversarial, competing} agents and the uncertainty of interacting with the natural world
- Historical reasons: For centuries humans have used them to exert their intelligence
- Fun, good, easy to understand
- Games often define very large search spaces
 - chess 35^{100} nodes in search tree, 10^{40} legal states

What Kinds of Games?

Mainly games of strategy with the following characteristics:

1. Sequence of **moves** to play
2. Rules that specify **possible moves**
3. Rules that specify a **payment** for each move
4. Objective is to **maximize** your payment

Games vs. Search Problems

- **Unpredictable opponent** → specifying a move for every possible opponent reply
- **Time limits** → unlikely to find goal, must approximate

Games

A game can be formally defined as a kind of search problem with 04 components:

1. The *initial state*, which includes the board position & identifies the player to move
2. A *successor function*, which returns a list of (move, state) pairs, each indicating a legal move & resulting state
3. A *terminal test*, which determines when the game is over. States where the game has ended are called terminal states
4. A *utility function*, which gives a numeric values for the terminal states.
 - In chess, the outcome is a win, loss, or draw, with values +1, -1, or 0.
 - Backgammon (+192~-192)

Game Trees

- Game trees are used to represent two-player games.
- Alternate moves in the game are represented by alternate levels in the tree.
- Nodes in the tree represent positions.
- Edges between nodes represent moves.
- Leaf nodes represent won, lost or drawn positions.

Game Tree (2-player, Deterministic, Turns)

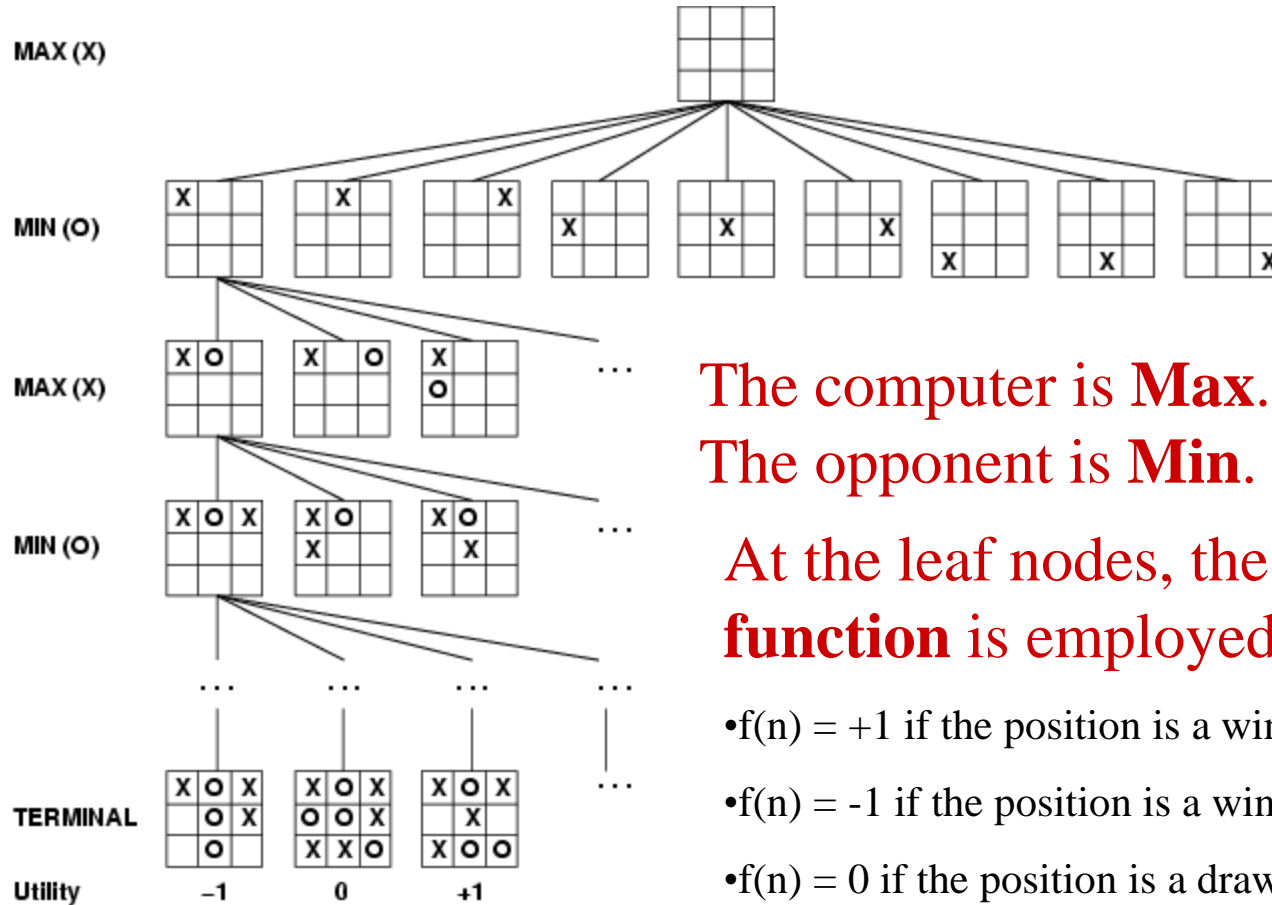
computer's
turn

opponent's
turn

computer's
turn

opponent's
turn

leaf nodes
are evaluated



The computer is **Max**.
The opponent is **Min**.

At the leaf nodes, the **utility function** is employed.

- $f(n) = +1$ if the position is a win for X.
- $f(n) = -1$ if the position is a win for O.
- $f(n) = 0$ if the position is a draw.

Assumptions

In talking about game playing systems, we make a number of assumptions:


- The opponent is rational – will play to win.
- The game is zero-sum – if one player wins, the other loses.
- Usually, the two players have complete knowledge of the game. For games such as poker, this is clearly not true.

Mini-Max Terminology

- ❑ **utility function:** the function applied to leaf nodes
- ❑ **backed-up value**
 - of a **max-position:** the value of its largest successor
 - of a **min-position:** the value of its smallest successor
- ❑ **minimax procedure:** search down several levels; at the bottom level apply the utility function, back-up values all the way up to the root node, and that node selects the move.

Minimax

- Perfect play for deterministic games
- Idea: choose move to position with highest **minimax value**
= best achievable payoff against best play
- E.g., 2-ply game:



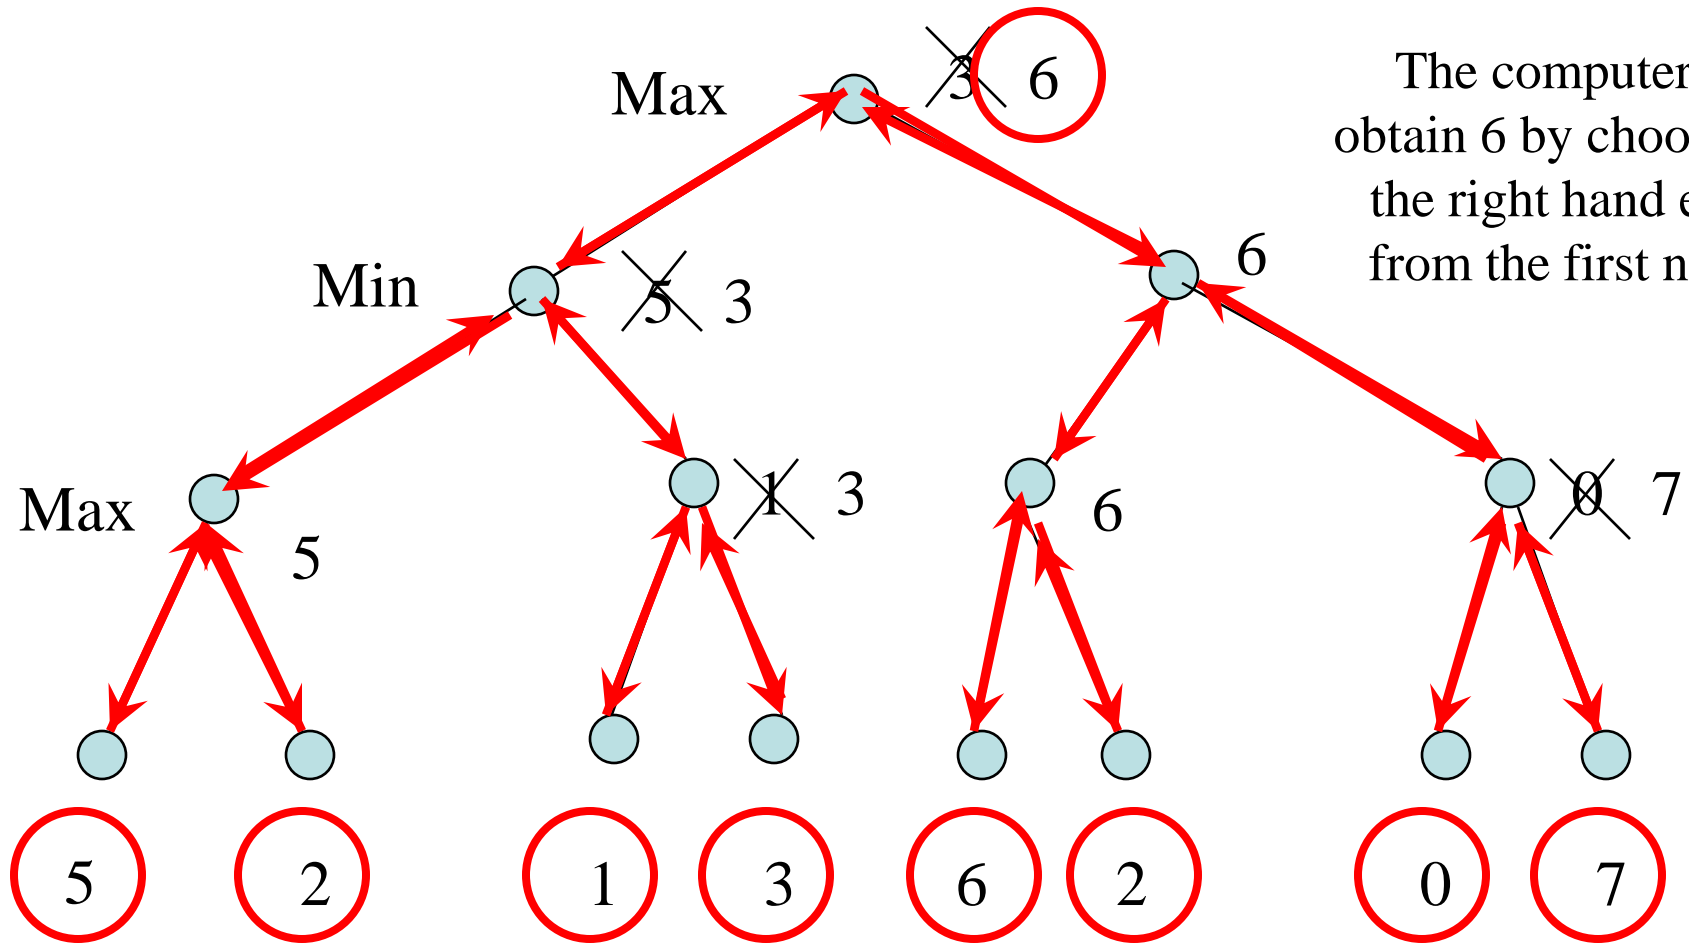
Minimax Strategy

- Why do we take the **min** value every other level of the tree?
 - These nodes represent the **opponent's** choice of move.
 - The computer assumes that the human will choose that move that is of **least value** to the computer.

Minimax Function

- $\text{MINIMAX-VALUE}(n) =$
 - $\text{UTILITY}(n)$ if n is a terminal state
 - $\max_{s \in \text{Successors}(n)} \text{MINIMAX-VALUE}(s)$ if n is a MAX node
 - $\min_{s \in \text{Successors}(n)} \text{MINIMAX-VALUE}(s)$ if n is a MIN node

Minimax – Animated Example



The computer can obtain 6 by choosing the right hand edge from the first node.

Minimax algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*


Properties of Minimax

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

Need to speed it up.

Alpha-Beta pruning

- We can improve on the performance of the minimax algorithm through **alpha-beta pruning**
- Basic idea: *“If you have an idea that is surely bad, don't take the time to see how truly awful it is.”* -- Pat Winston




- We don't need to compute the value at this node.
- No matter what it is, it can't affect the value of the root node.


Alpha-Beta pruning

- Traverse the search tree in depth-first order
- At each **MAX** node n , **alpha**(n) = maximum value found so far
- At each **MIN** node n , **beta**(n) = minimum value found so far
 - The alpha values start at $-\infty$ and only increase, while beta values start at $+\infty$ and only decrease
- **Beta cutoff:** Given MAX node n , cut off search below n (i.e., don't generate/examine any more of n 's children) if $\text{alpha}(n) \geq \text{beta}(i)$ for some MIN node ancestor i of n .
- **Alpha cutoff:** stop searching below MIN node n if $\text{beta}(n) \leq \text{alpha}(i)$ for some MAX node ancestor i of n .


Alpha-Beta General example




Alpha-Beta Pruning Example-2



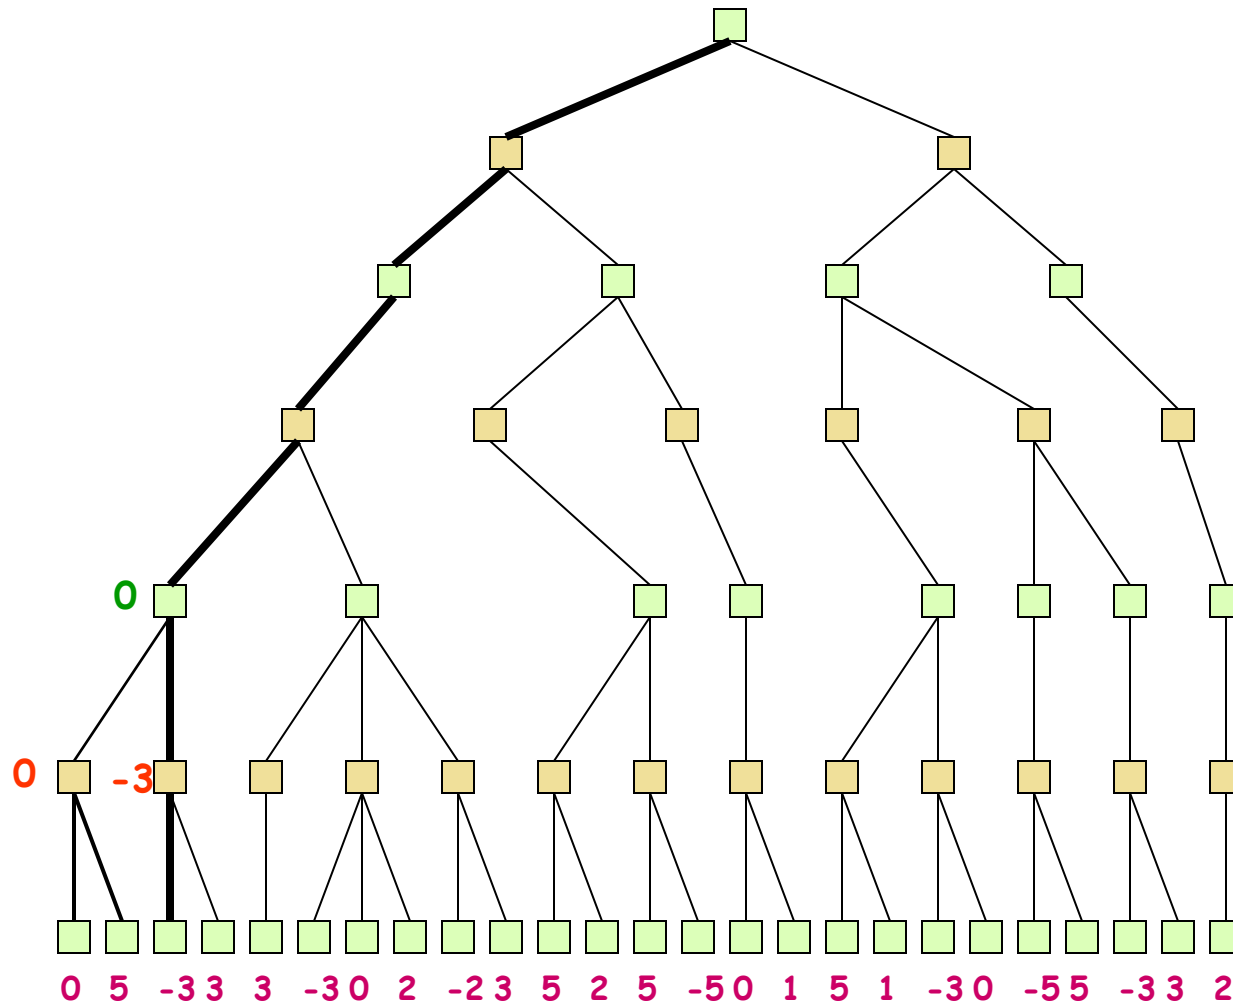
Alpha-Beta Pruning Example-2



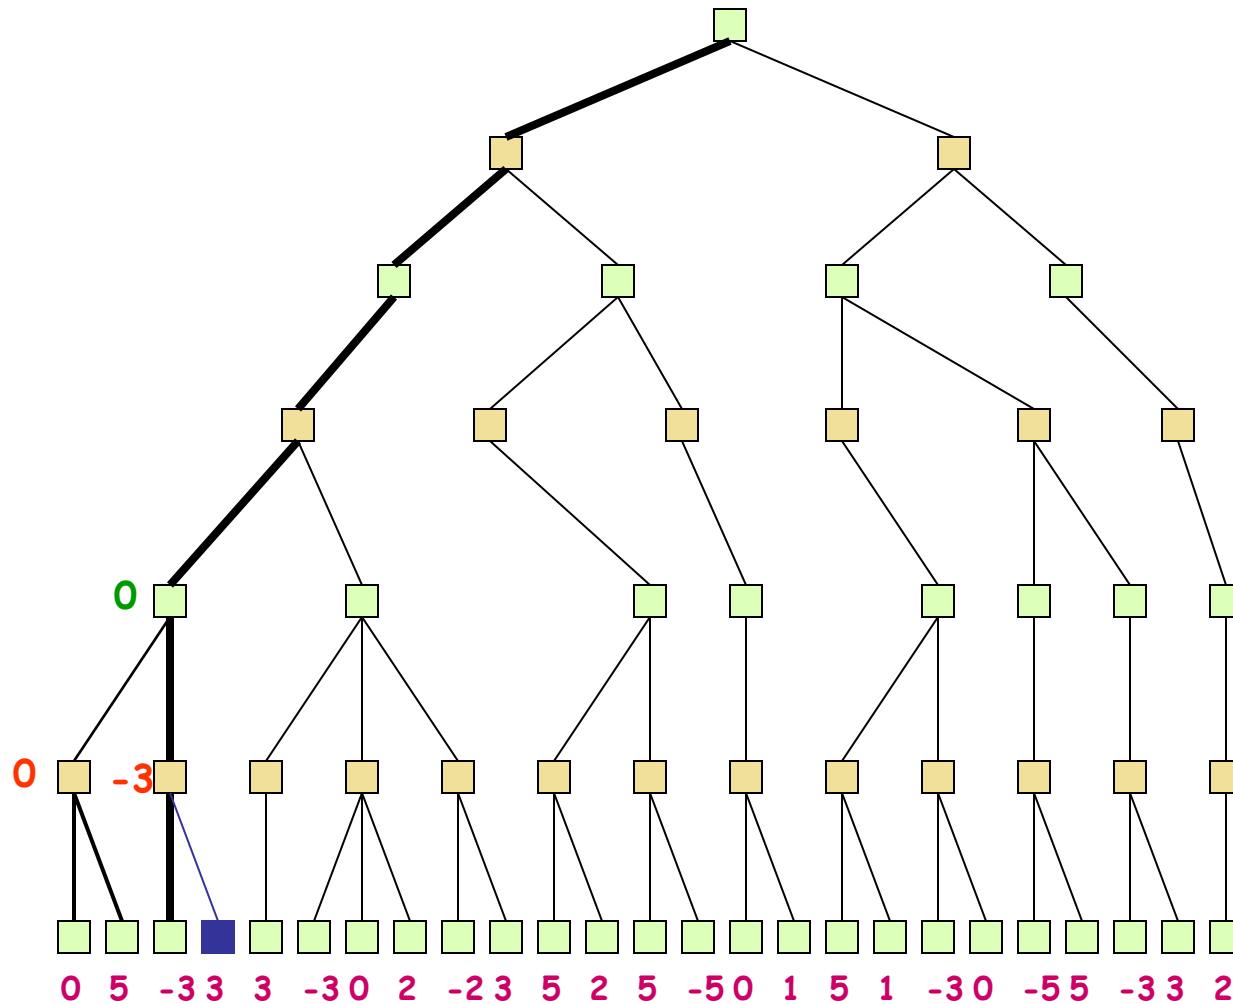
Alpha-Beta Pruning Example-2



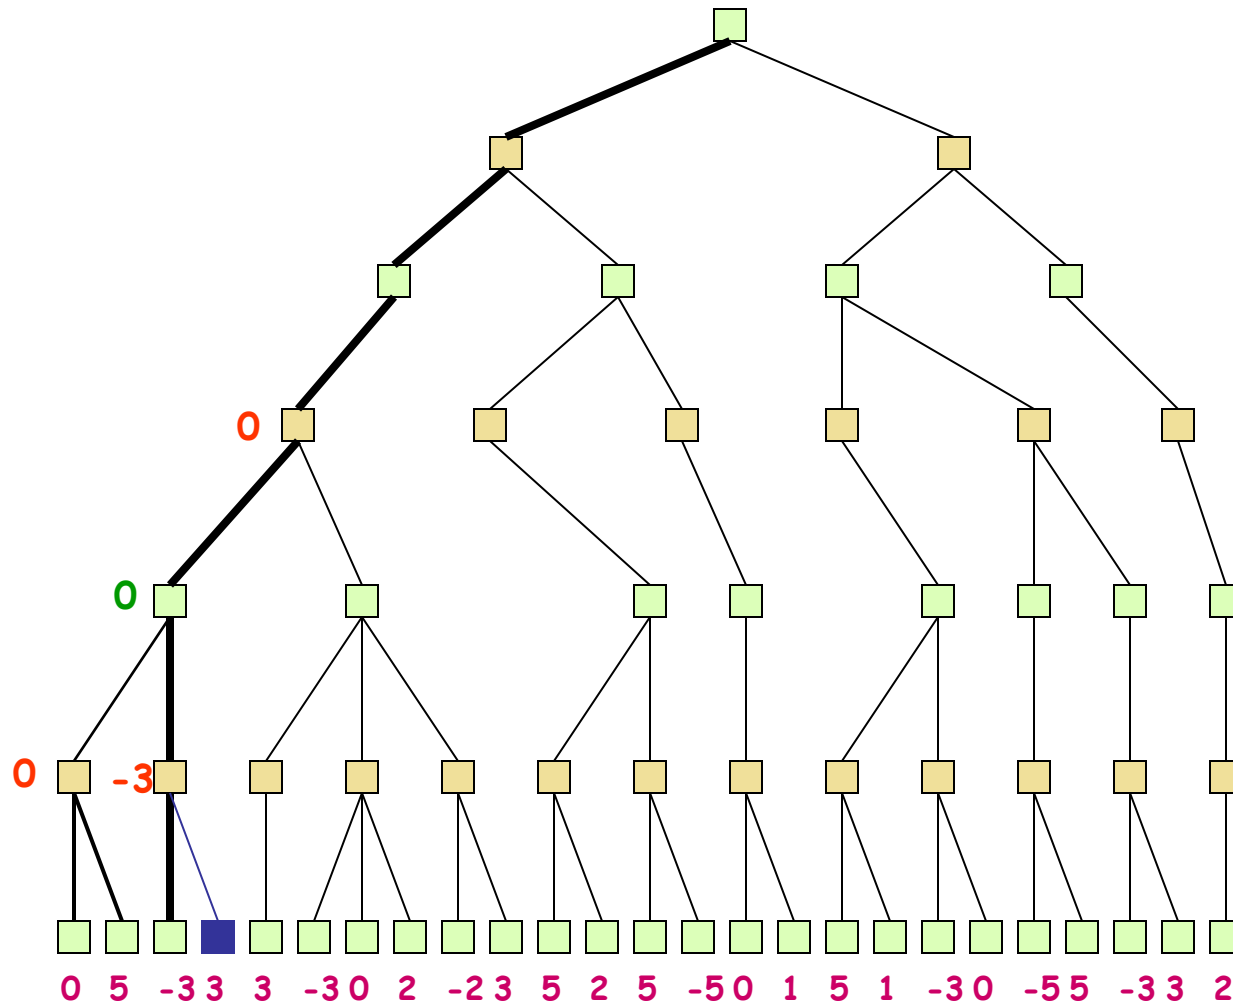
Alpha-Beta Pruning Example-2



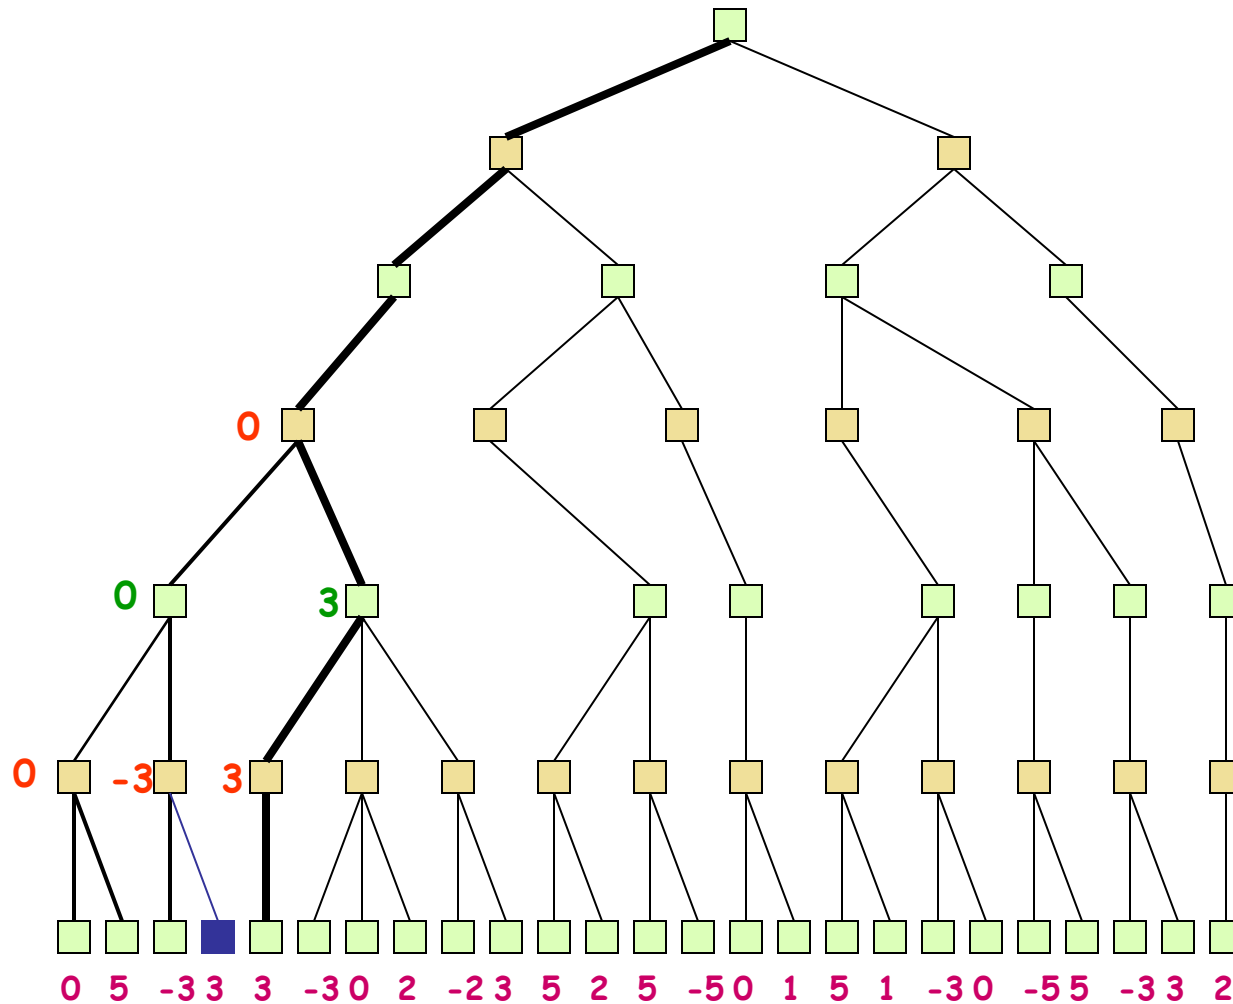
Alpha-Beta Pruning Example-2



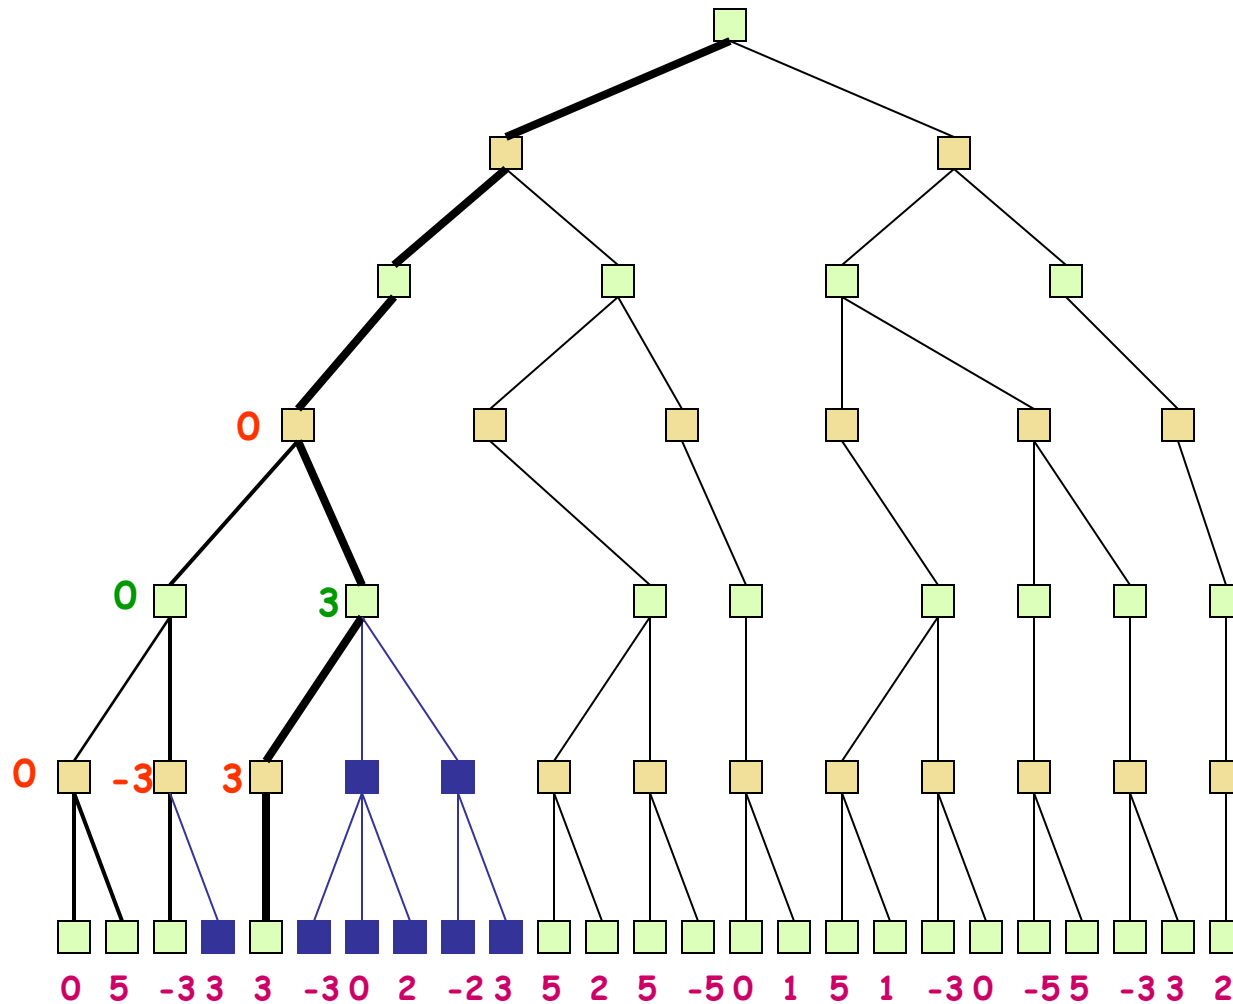
Alpha-Beta Pruning Example-2



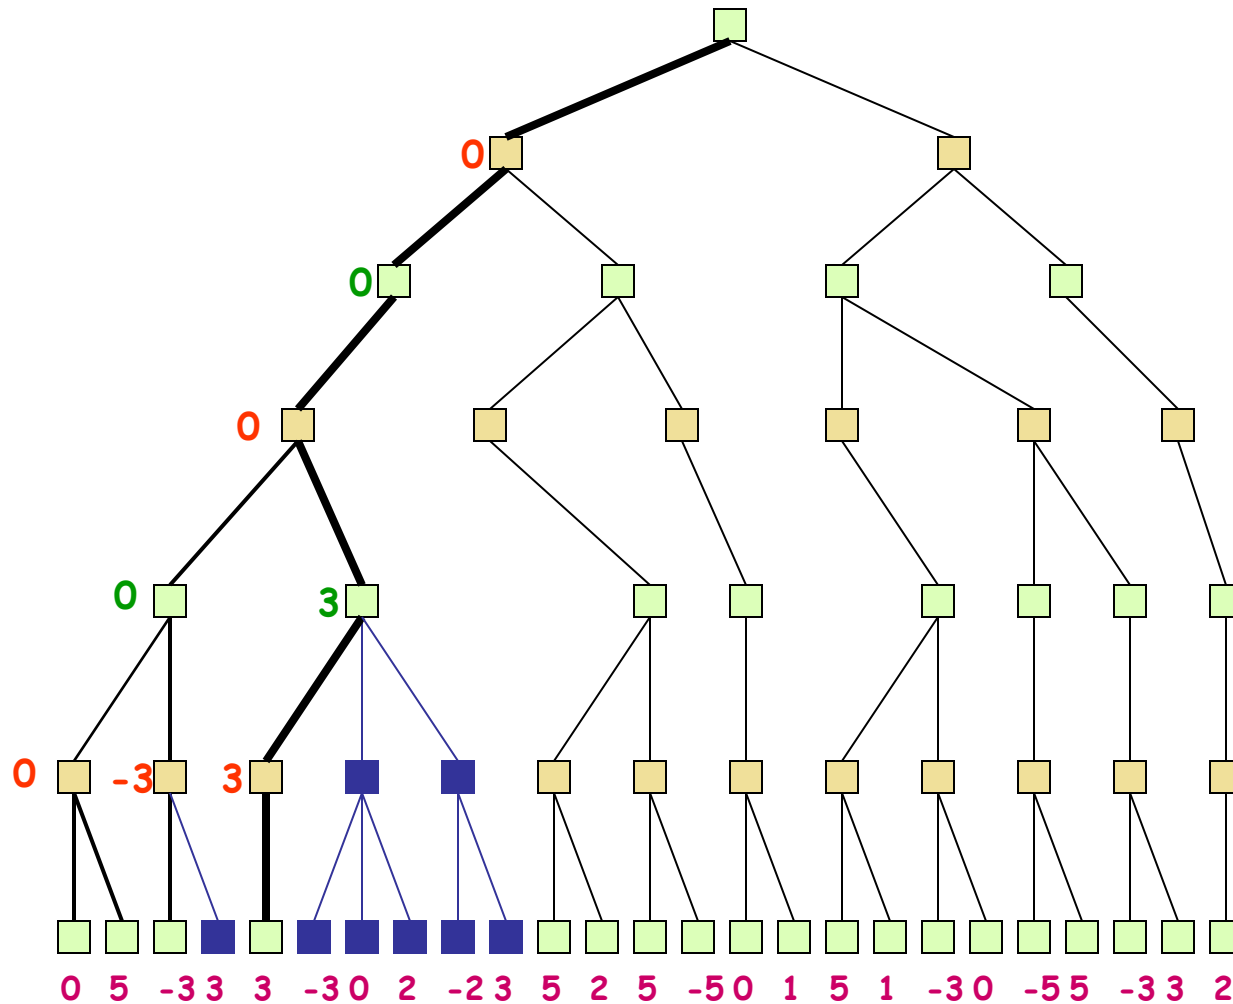
Alpha-Beta Pruning Example-2



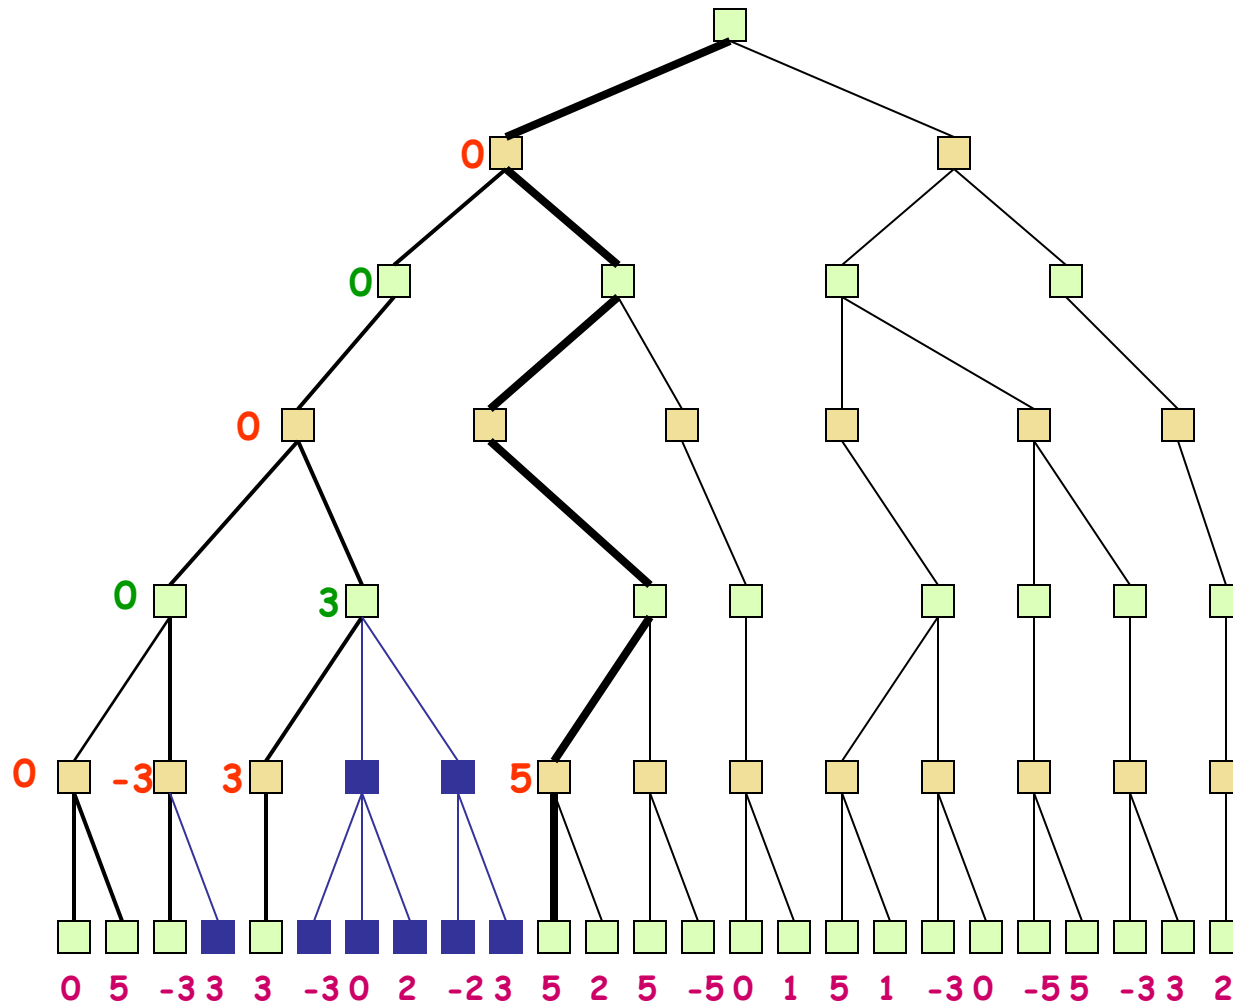
Alpha-Beta Pruning Example-2



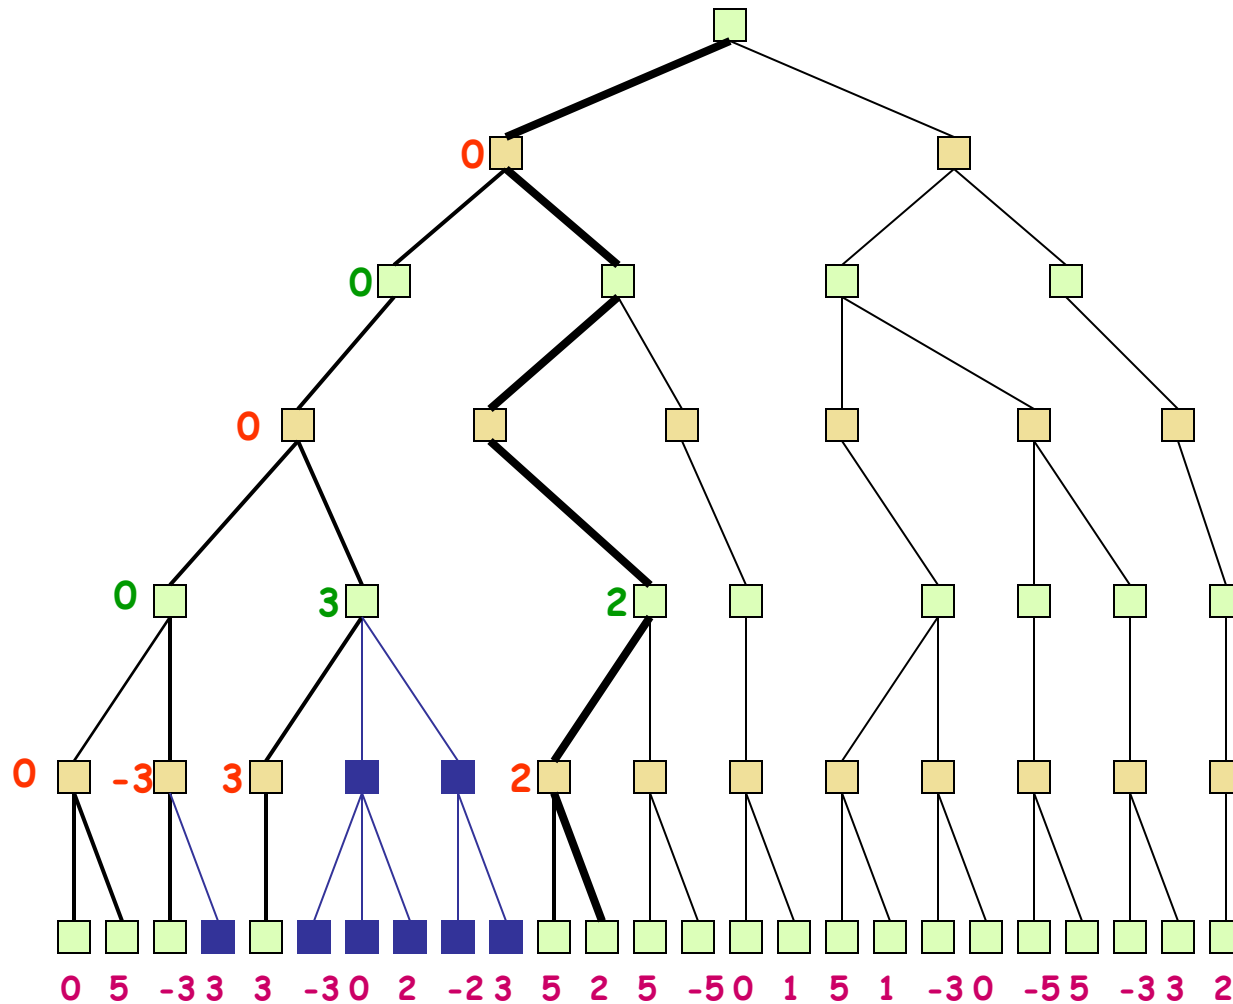
Alpha-Beta Pruning Example-2



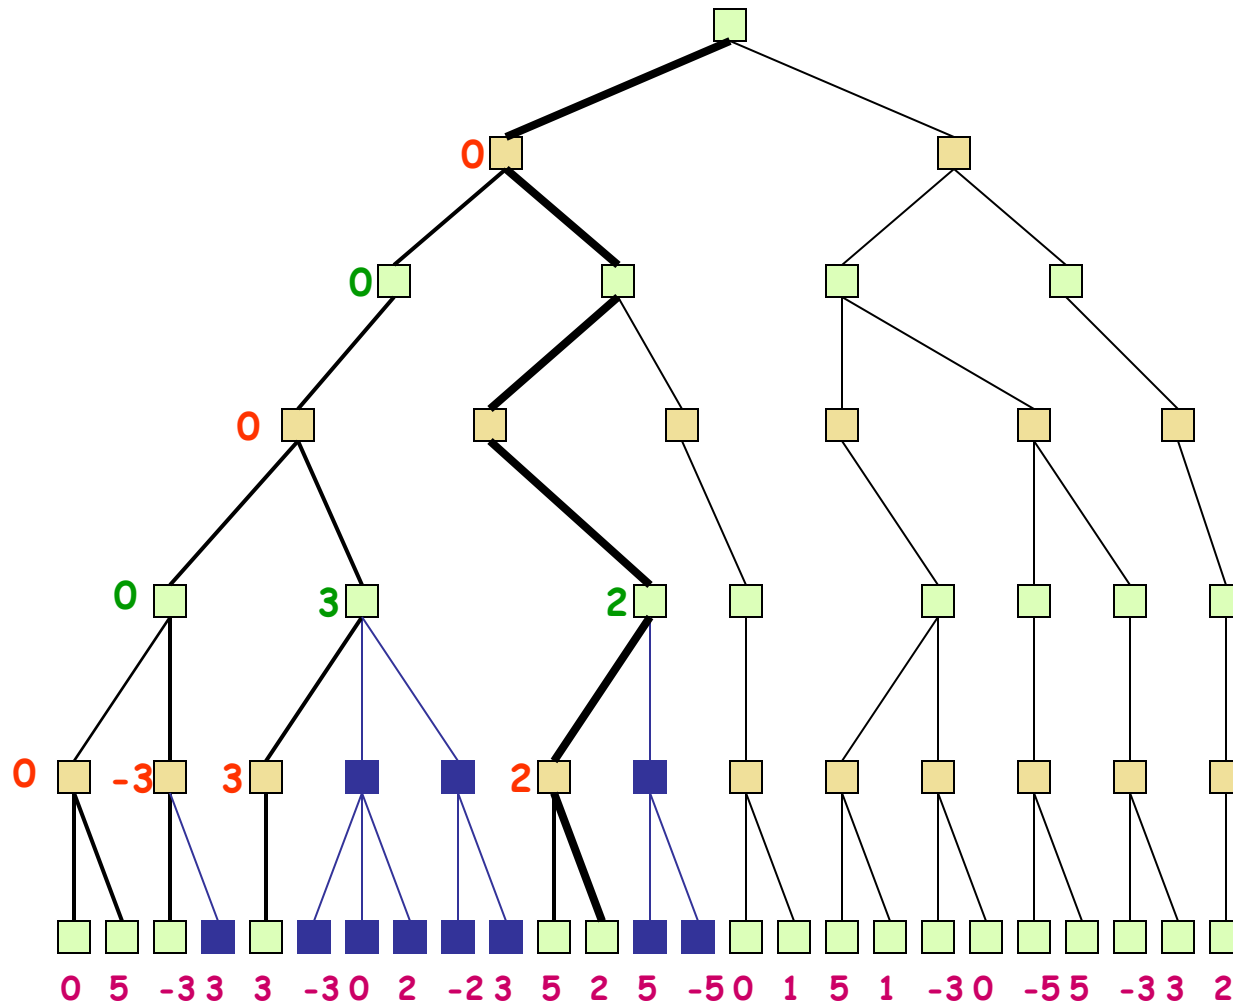
Alpha-Beta Pruning Example-2



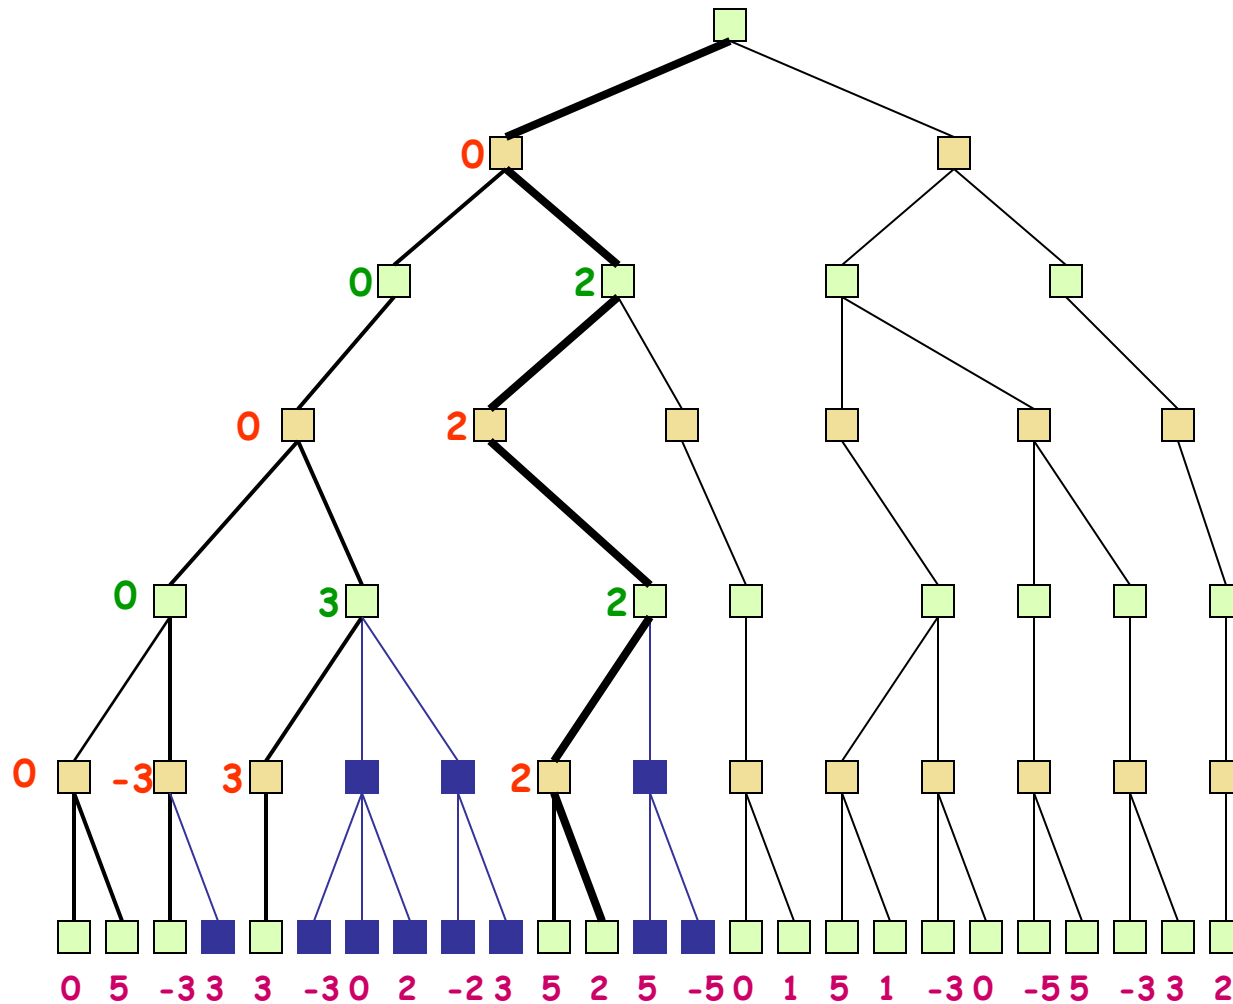
Alpha-Beta Pruning Example-2



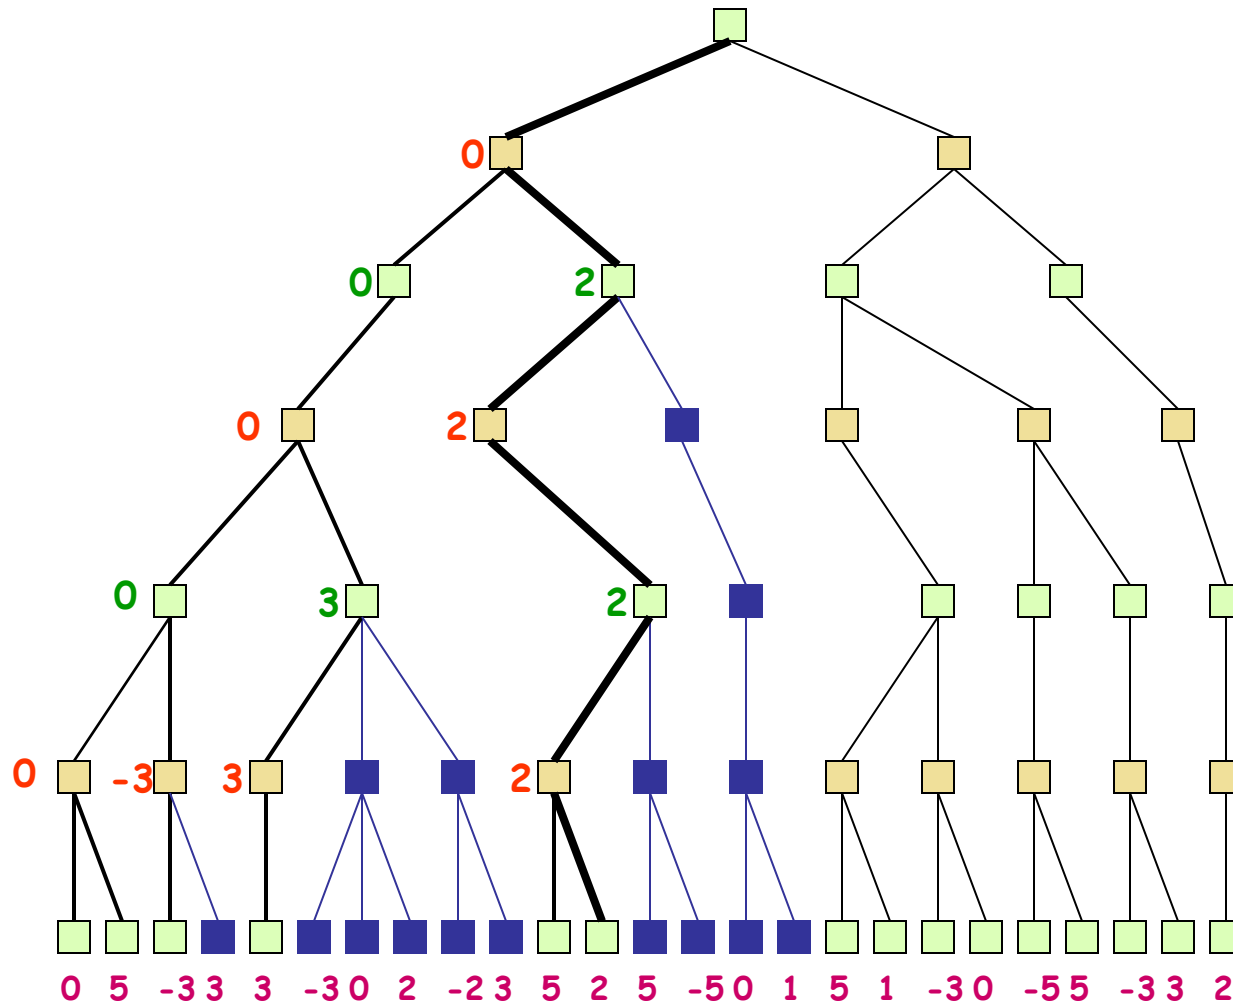
Alpha-Beta Pruning Example-2



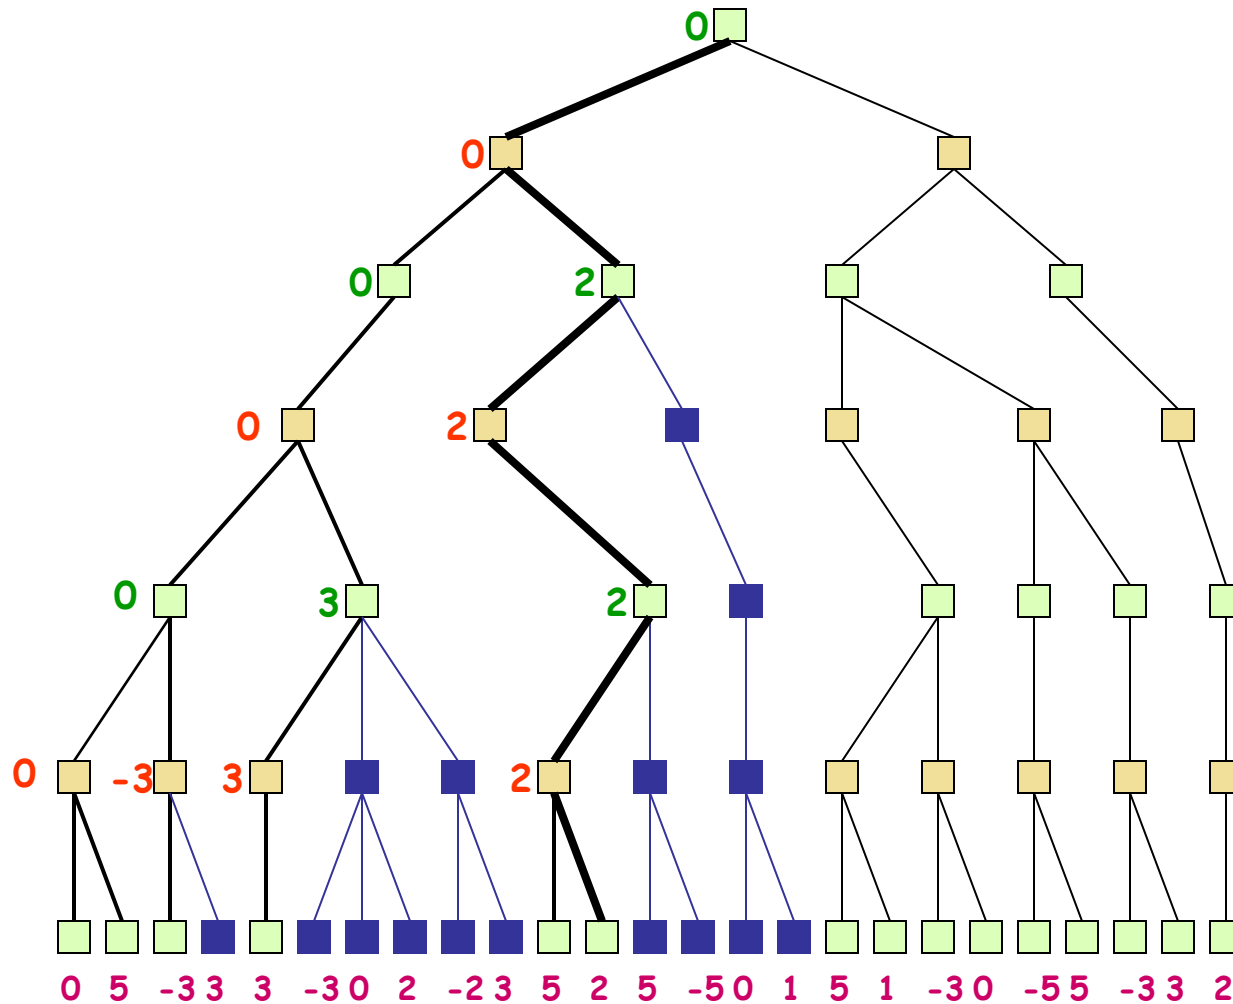
Alpha-Beta Pruning Example-2



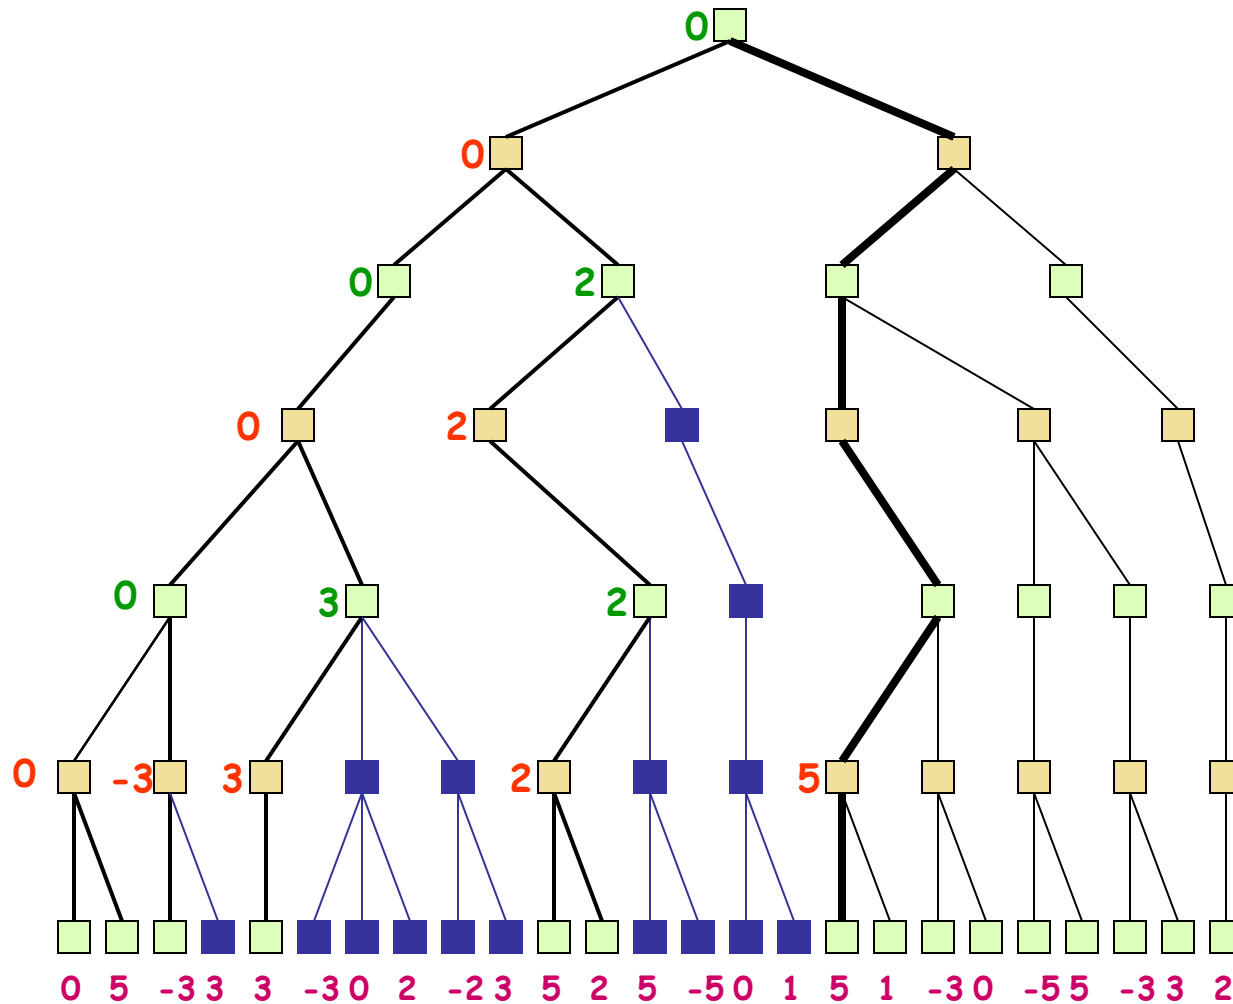
Alpha-Beta Pruning Example-2



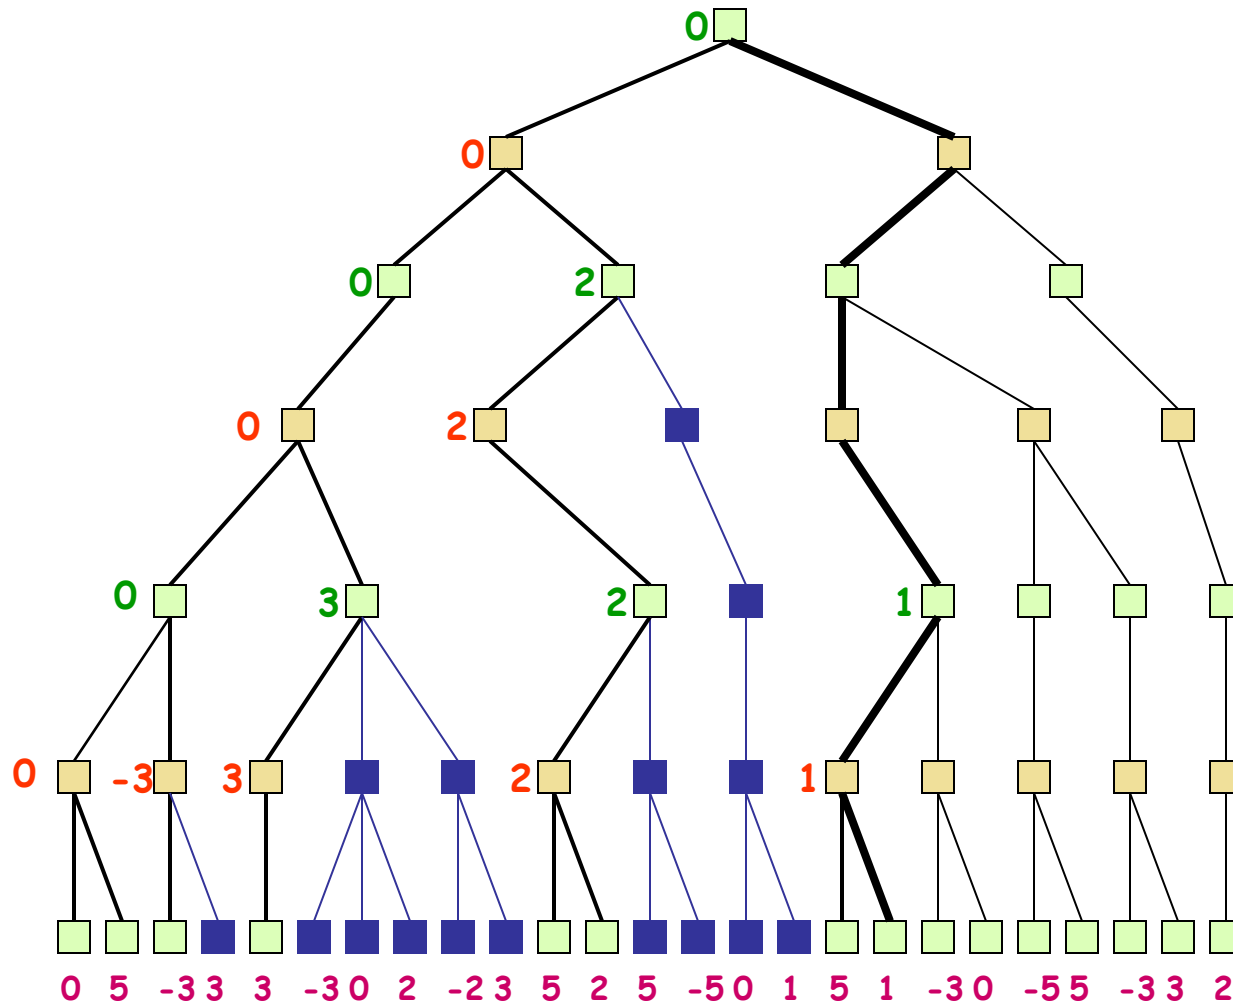
Alpha-Beta Pruning Example-2



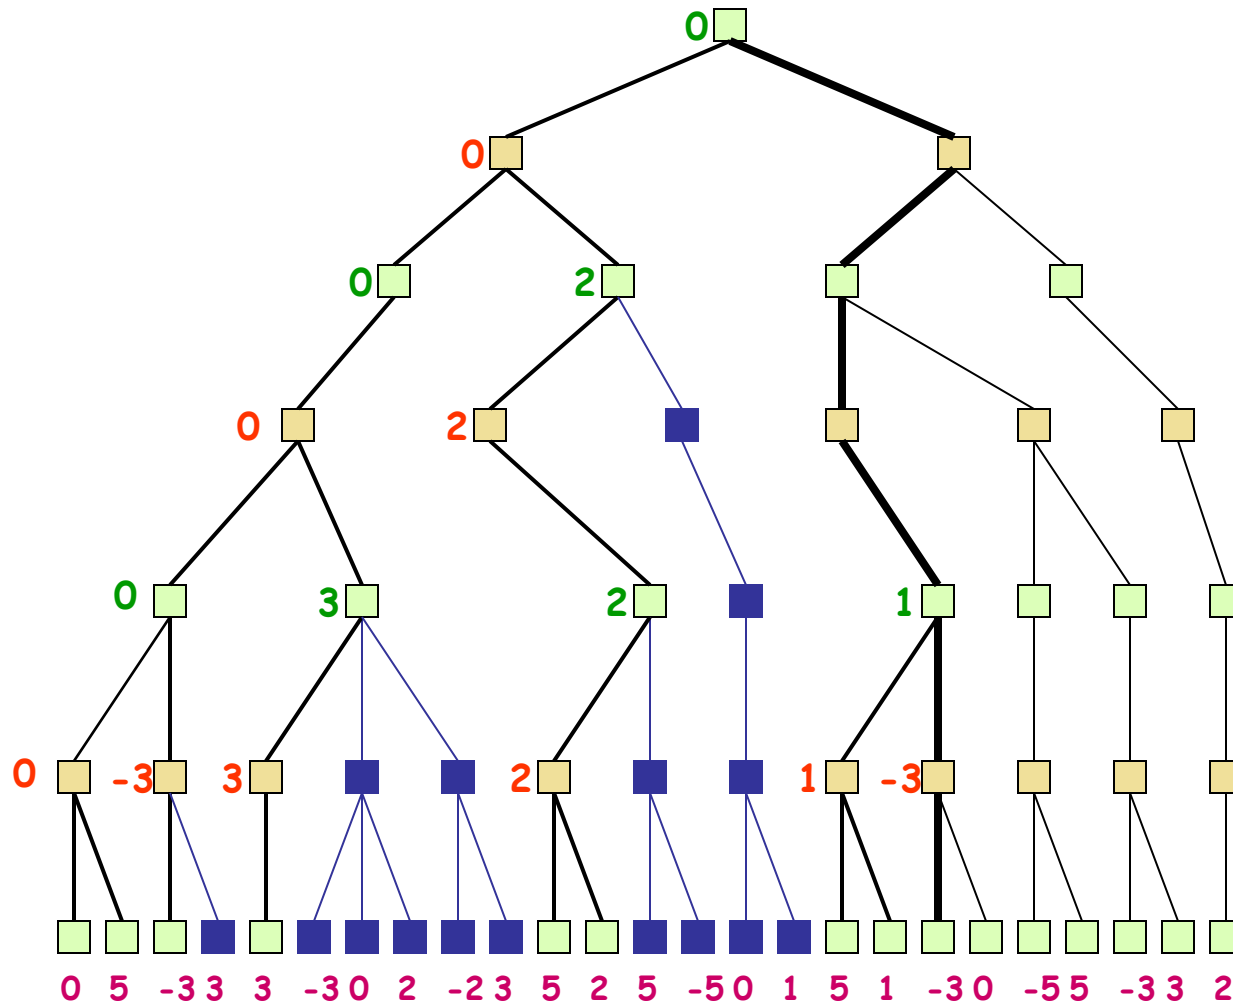
Alpha-Beta Pruning Example-2



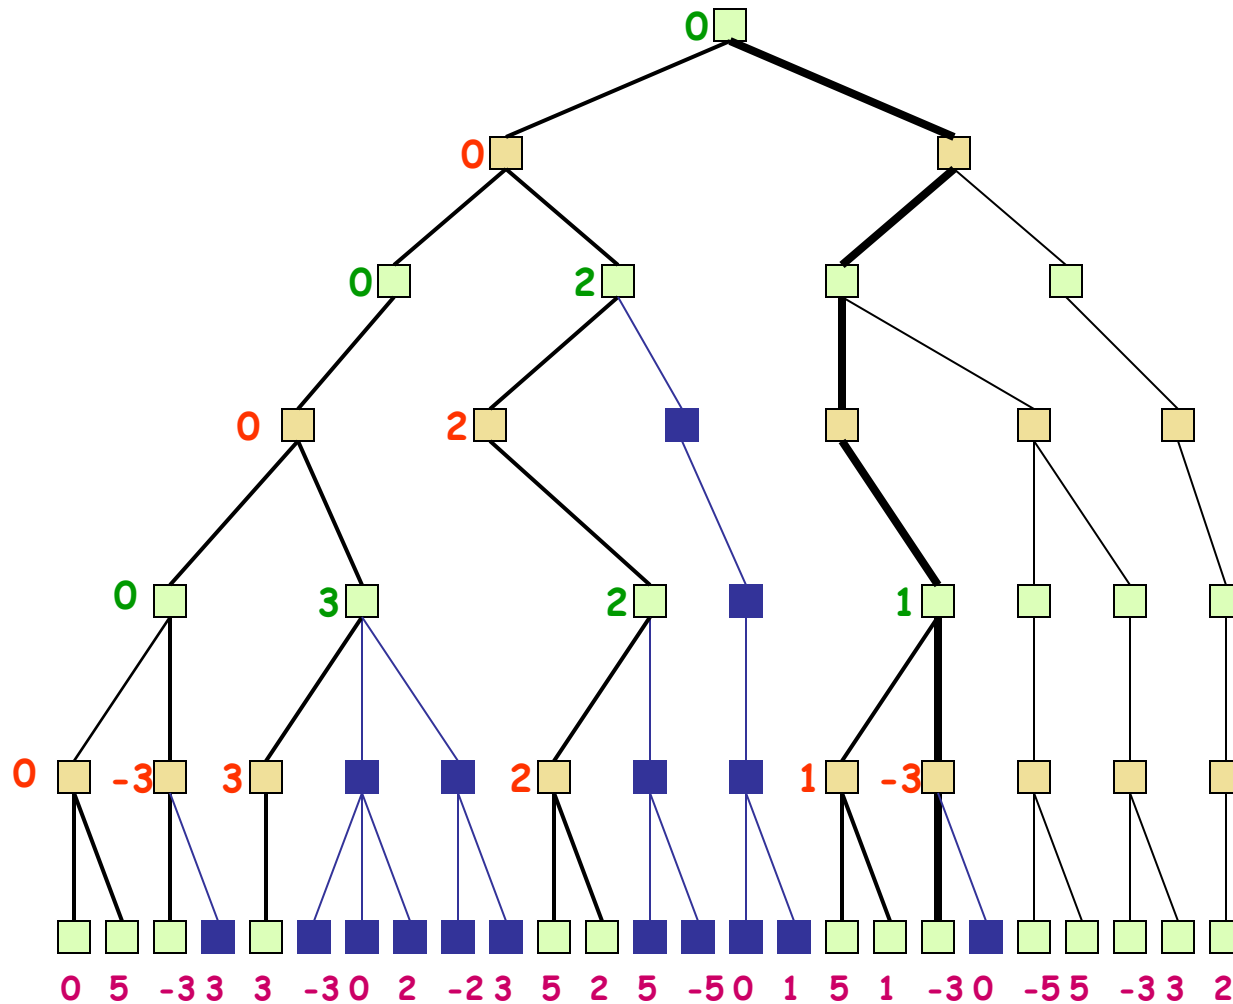
Alpha-Beta Pruning Example-2



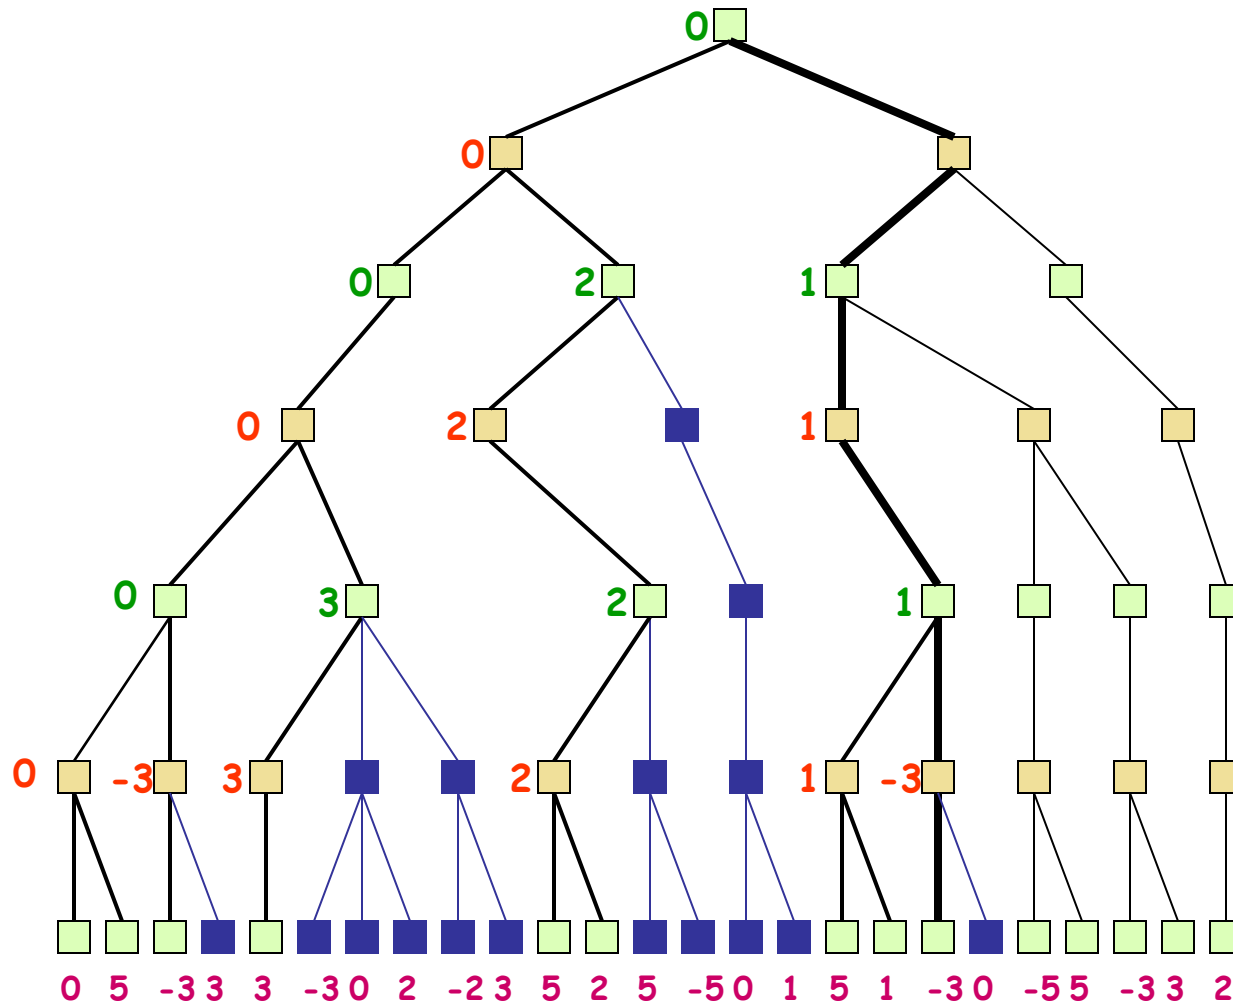
Alpha-Beta Pruning Example-2



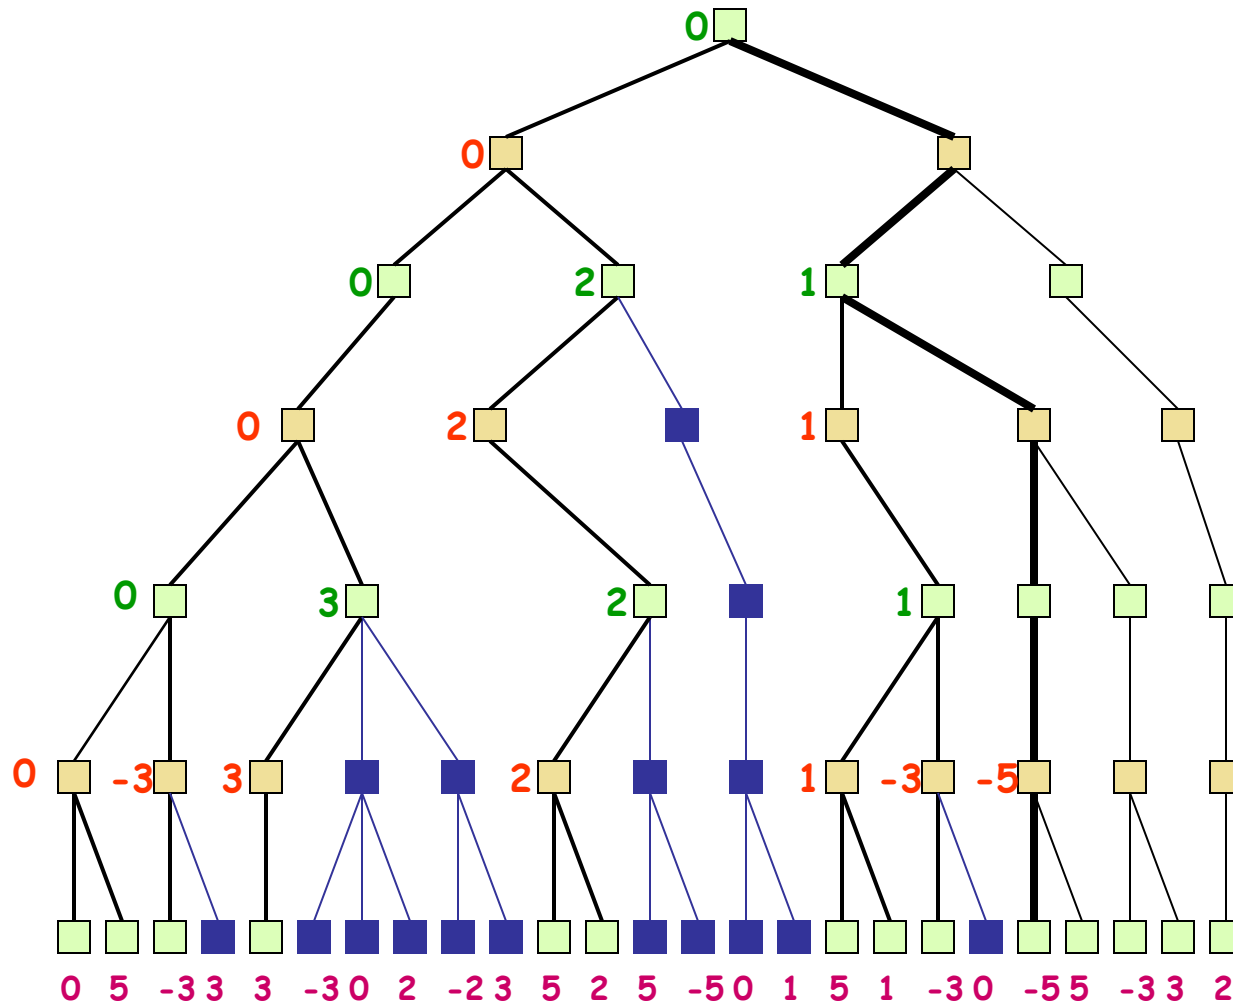
Alpha-Beta Pruning Example-2



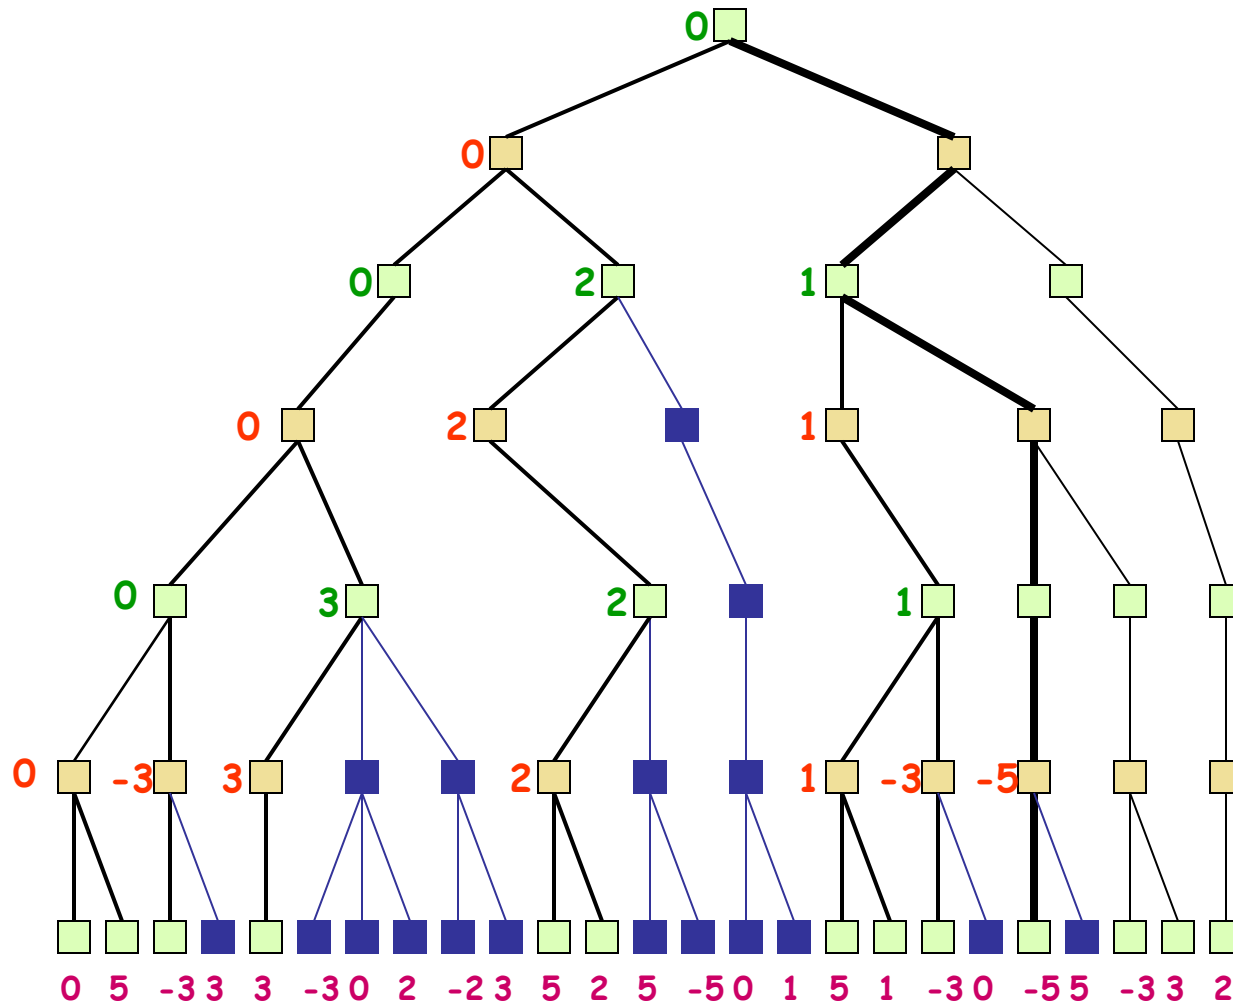
Alpha-Beta Pruning Example-2



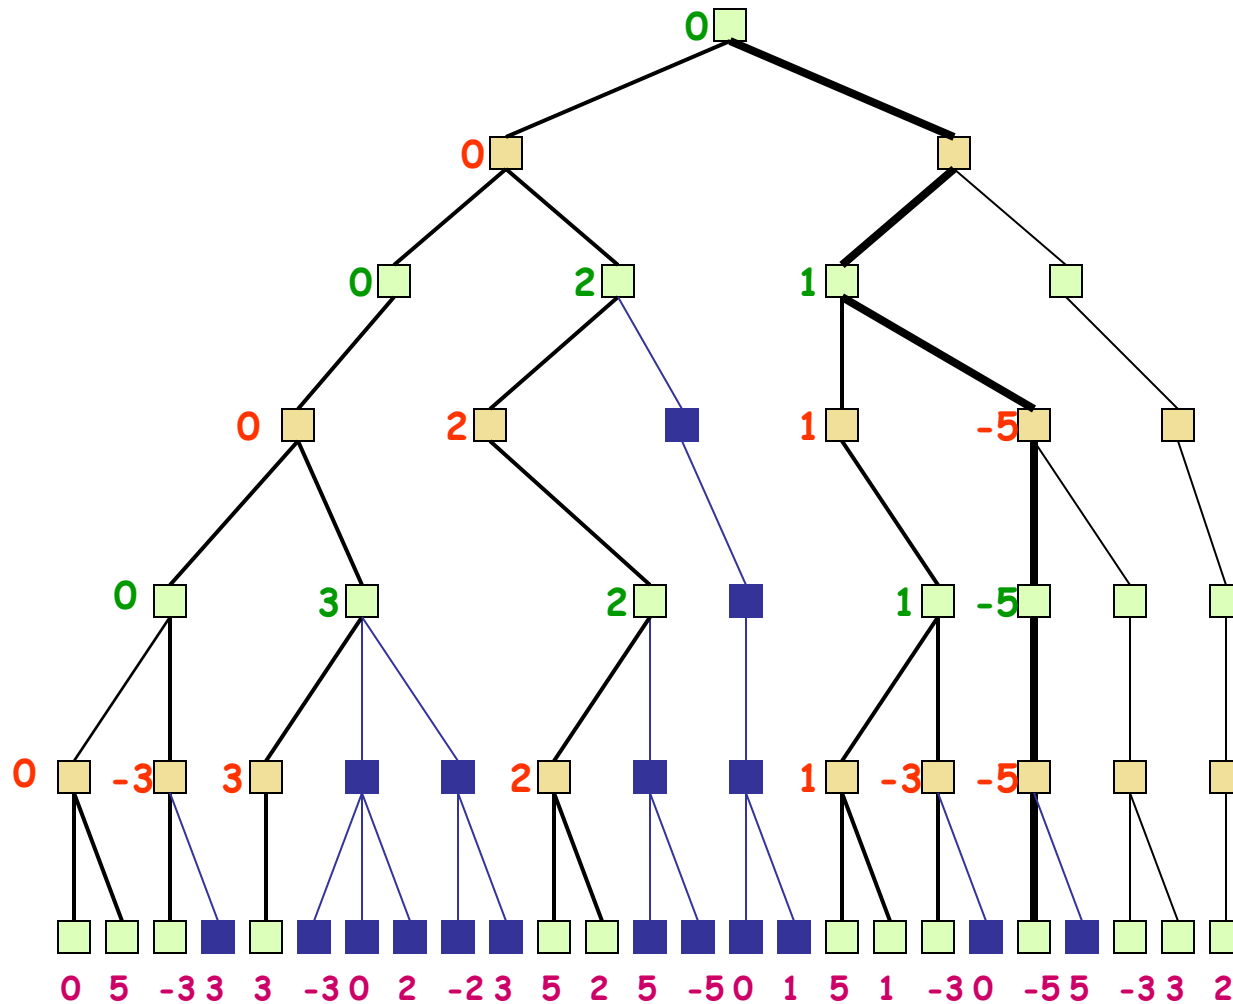
Alpha-Beta Pruning Example-2



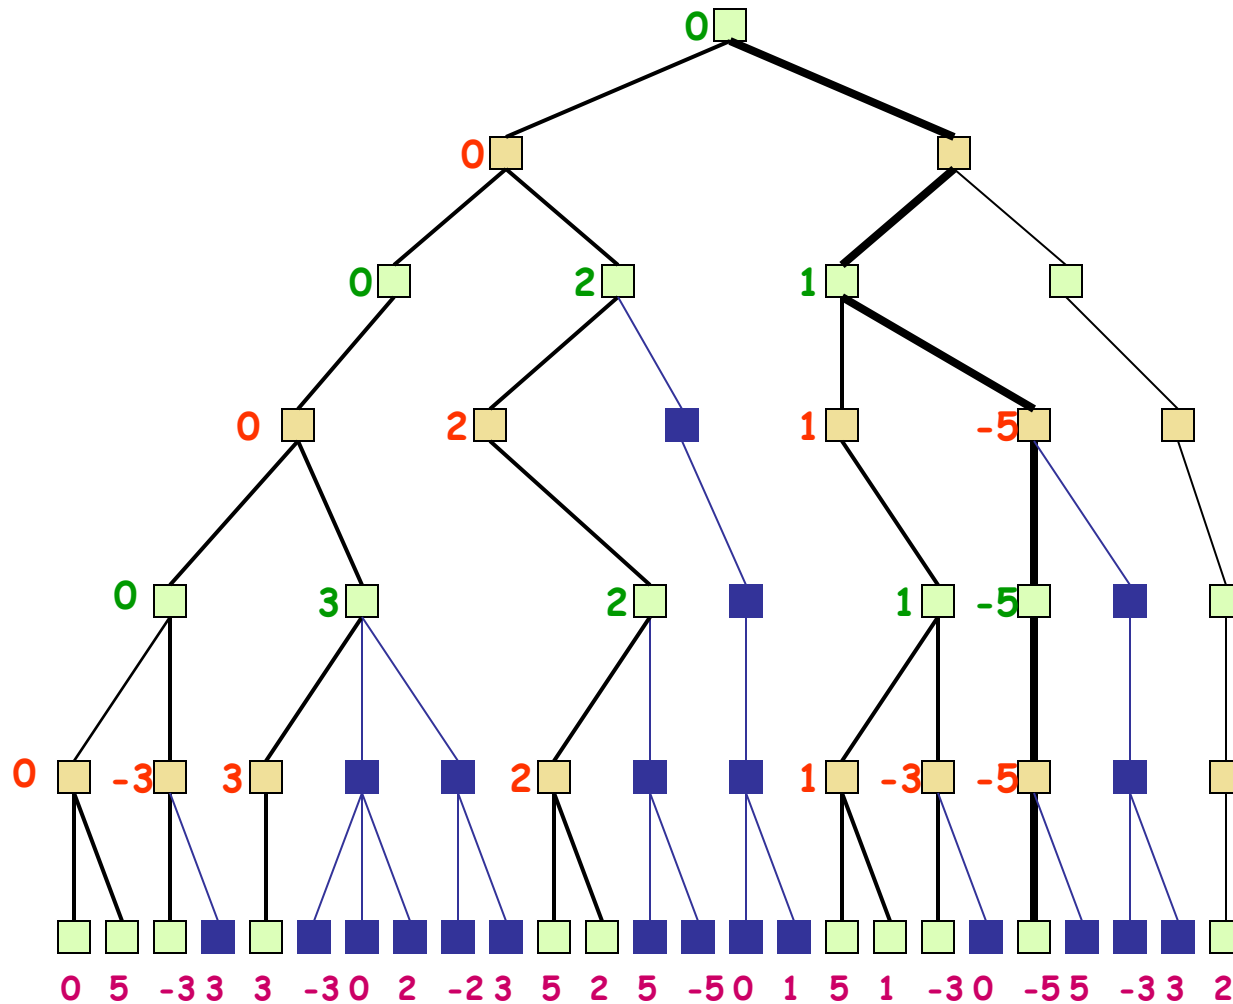
Alpha-Beta Pruning Example-2



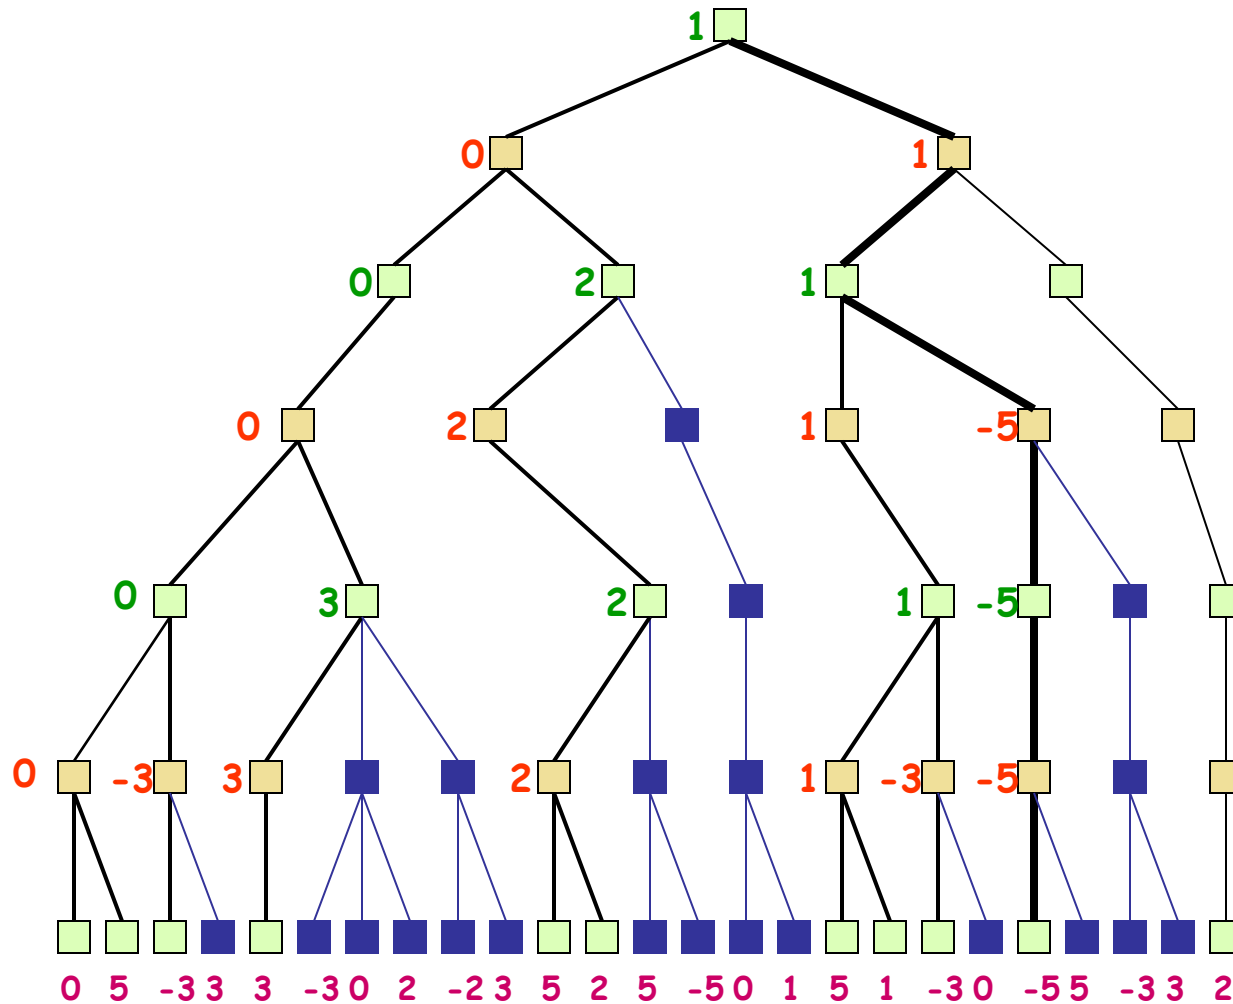
Alpha-Beta Pruning Example-2



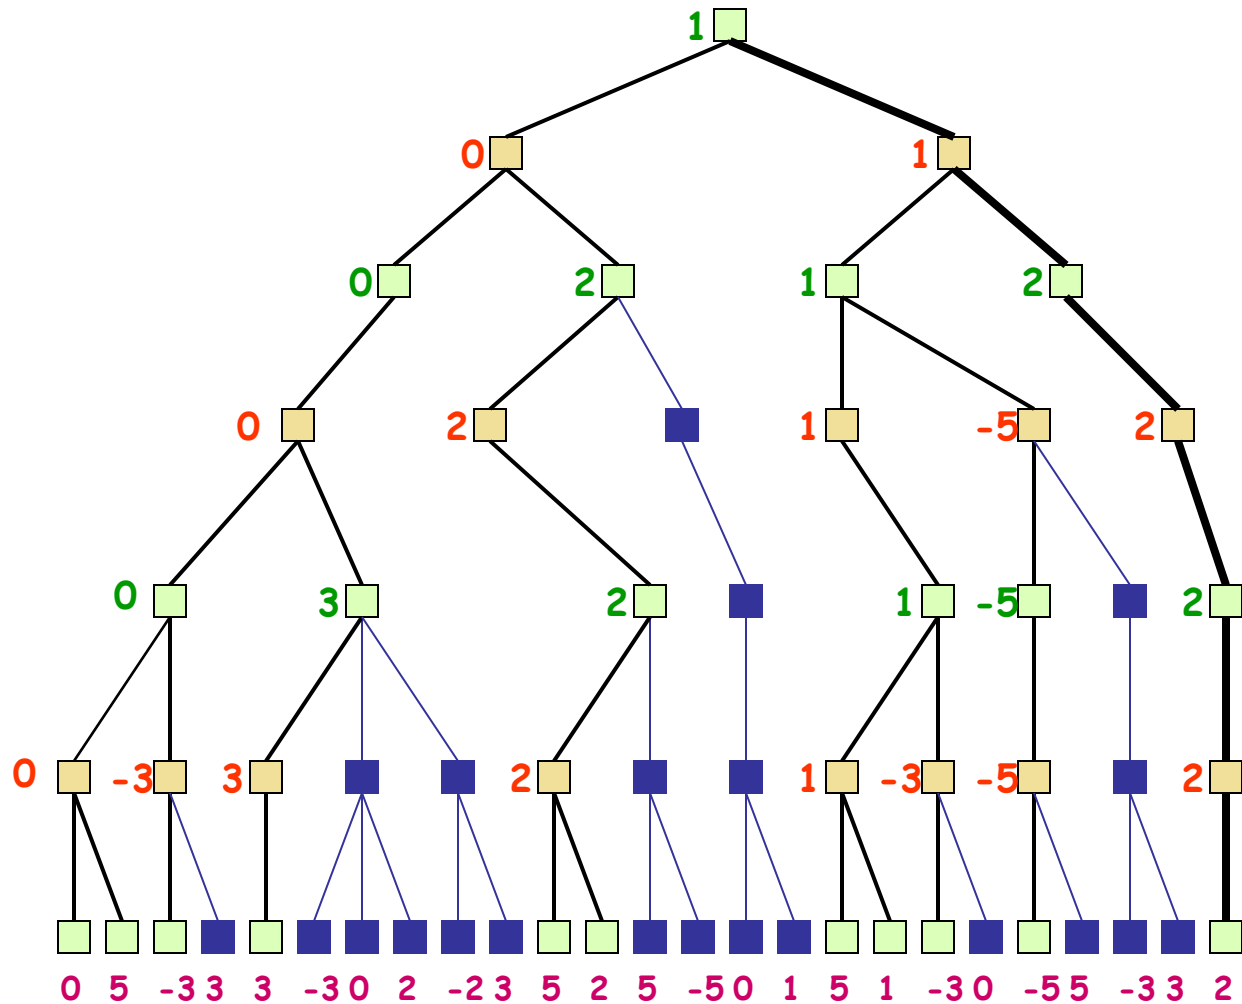
Alpha-Beta Pruning Example-2



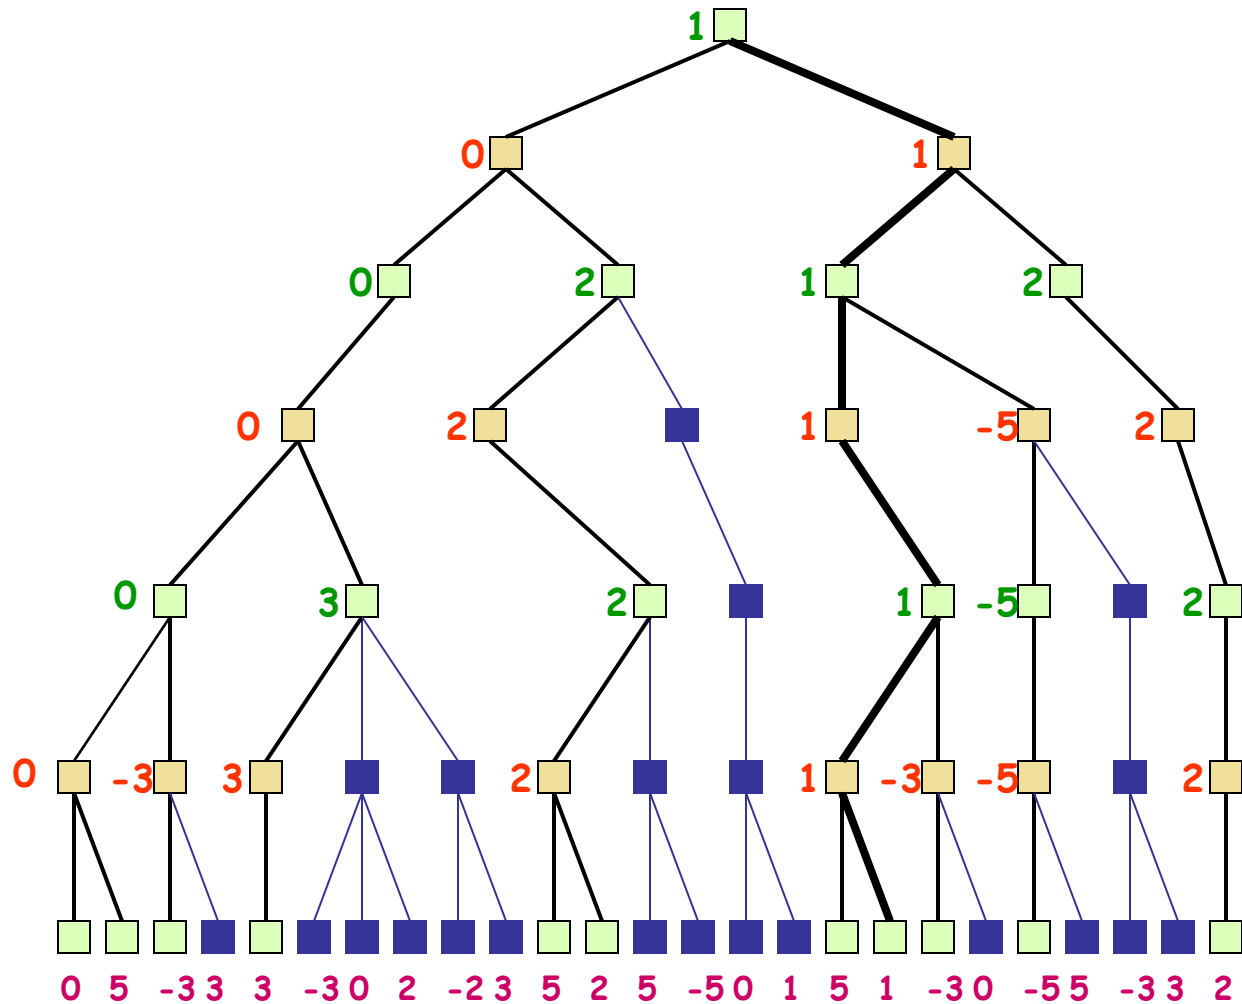
Alpha-Beta Pruning Example-2



Alpha-Beta Pruning Example-2



Alpha-Beta Pruning Example-2



The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$


for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v **cutoff**

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return v



The α - β algorithm

function MIN-VALUE($state, \alpha, \beta$) *returns a utility value*

inputs: $state$, current state in game

α , the value of the best alternative for MAX along the path to $state$

β , the value of the best alternative for MIN along the path to $state$

if TERMINAL-TEST($state$) **then return** UTILITY($state$)

$v \leftarrow +\infty$


for a, s in SUCCESSORS($state$) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** v **cutoff**

$\beta \leftarrow \text{MIN}(\beta, v)$

return v




Why is it called α - β ?

□ α is the value of the best (i.e., highest-value) choice found so far for *MAX* at any choice point along the path to the root.

- If v is worse than α , *MAX* will avoid it
→ prune that branch

□ β is the value of the best (i.e., lowest-value) choice found so far for *MIN* at any choice point along the path for to the root.



Properties of α - β

- Pruning **does not** affect final result. This means that it **gets the exact same result as does full minimax**.
- Good move ordering improves effectiveness of pruning
- With "perfect ordering," time complexity = $O(b^{m/2})$
→ **doubles** depth of search
- A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)

State of the art

- **Chess:**
 - Deep Blue beat Gary Kasparov in 1997
 - Garry Kasparov vs. Deep Junior (Feb 2003): tie!
 - Kasparov vs. X3D Fritz (November 2003): tie!
- **Checkers:** Chinook is the world champion
- **Go:** Computer players are decent, at best
- **Bridge:** “Expert” computer players exist, but no world champions yet
- **Poker:** CPRG regularly beats human experts
- Check out: <http://www.cs.ualberta.ca/~games/>

Summary

- Games are fun to work on!
- They illustrate several important points about AI.
- Perfection is unattainable → must approximate.
- Game playing programs have shown the world what AI can do.